

Practical 2

Emilia Löscher

19-9-2022

First, load the packages:

```
library(ISLR)
```

```
## Warning: Paket 'ISLR' wurde unter R Version 4.1.3 erstellt
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v ggplot2 3.3.5      v purrr   0.3.4
## v tibble  3.1.6      v dplyr  1.0.8
## v tidyr   1.2.0      v stringr 1.4.0
## v readr   2.1.2      v forcats 0.5.1
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

Inspect the data

```
head(Hitters)
```

```
##           AtBat Hits HmRun Runs RBI Walks Years CATBat CHits CHmRun
## -Andy Allanson    293   66     1   30  29   14     1    293    66     1
## -Alan Ashby       315   81     7   24  38   39   14   3449   835    69
## -Alvin Davis      479  130    18   66  72   76    3   1624   457    63
```

## -Andre Dawson	496	141	20	65	78	37	11	5628	1575	225
## -Andres Galarrraga	321	87	10	39	42	30	2	396	101	12
## -Alfredo Griffin	594	169	4	74	51	35	11	4408	1133	19
##	CRuns	CRBI	CWalks	League	Division	PutOuts	Assists	Errors		
## -Andy Allanson	30	29	14	A	E	446	33	20		
## -Alan Ashby	321	414	375	N	W	632	43	10		
## -Alvin Davis	224	266	263	A	W	880	82	14		
## -Andre Dawson	828	838	354	N	E	200	11	3		
## -Andres Galarrraga	48	46	33	N	E	805	40	4		
## -Alfredo Griffin	501	336	194	A	W	282	421	25		
##	Salary	NewLeague								
## -Andy Allanson	NA	A								
## -Alan Ashby	475.0	N								
## -Alvin Davis	480.0	A								
## -Andre Dawson	500.0	N								
## -Andres Galarrraga	91.5	N								
## -Alfredo Griffin	750.0	A								


1

Name the aesthetics, geoms, scales, and facets of the above visualisation. Also name any statistical transformations or special coordinate systems.

```
homeruns_plot <-
  ggplot(Hitters, aes(x = Hits, y = HmRun)) +
  geom_point() +
  labs(x = "Hits", y = "Home runs")

homeruns_plot +
  geom_density_2d() +
  labs(title = "Cool density and scatter plot of baseball data") +
  theme_minimal()
```

Cool density and scatter plot of baseball data



- Aesthetics: x-axis is Hits, y-axis is HomeRun [aes(x = Hits, y = HmRun)]
- Geoms: scatter plot + density lines
- Scales: x: 0-250, y: 0-42 (both continuous)
- Facets: -
- Statistical transformation: -
- Special Coordinate System: minimal theme

2

Run the code below to generate data. There will be three vectors in your environment. Put them in a data frame for entering it in a ggplot() call using either the data.frame() or the tibble() function. Give informative names and make sure the types are correct (use the as.() functions). Name the result gg_students

```
set.seed(1234)
student_grade <- rnorm(32, 7)
```

```

student_number <- round(runif(32) * 2e6 + 5e6)
programme      <- sample(c("Science", "Social Science"), 32, replace = TRUE)

gg_students <- data.frame(as.numeric(student_grade), as.character(student_number), as.factor(programme))

colnames(gg_students) <- c("Grade", "Student number", "Programme")

gg_students


```

##	Grade	Student number	Programme
## 1	5.792934	5478051	Social Science
## 2	7.277429	6412989	Science
## 3	8.084441	5616190	Social Science
## 4	4.654302	6017095	Social Science
## 5	7.429125	5103293	Social Science
## 6	7.506056	6129140	Science
## 7	6.425260	5242960	Science
## 8	6.453368	6785673	Science
## 9	6.435548	5029255	Science
## 10	6.109962	6566242	Social Science
## 11	6.522807	5179923	Science
## 12	6.001614	6038380	Social Science
## 13	6.223746	5768533	Science
## 14	7.064459	5140105	Science
## 15	7.959494	5641289	Science
## 16	6.889715	6336991	Social Science
## 17	6.488990	6852801	Social Science
## 18	6.088805	5943819	Social Science
## 19	6.162828	5285231	Social Science
## 20	9.415835	6088540	Science
## 21	7.134088	5392349	Social Science
## 22	6.509314	6797161	Science
## 23	6.559452	5779000	Social Science
## 24	7.459589	5621742	Science
## 25	6.306280	5320057	Science
## 26	5.551795	6792372	Social Science
## 27	7.574756	5332788	Social Science
## 28	5.976344	6800849	Science
## 29	6.984862	5268156	Social Science
## 30	6.064051	5263228	Science
## 31	8.102298	5210575	Social Science
## 32	6.524407	6023167	Social Science

3

Plot the first homeruns_plot again, but map the Hits to the y-axis and the HmRun to the x-axis instead.

```
homeruns_plot_3 <-  
  ggplot(Hitters, aes(x = HmRun, y = Hits)) +  
  geom_point() +  
  labs(y = "Hits", x = "Home runs")  
  
homeruns_plot_3
```




4

Recreate the same plot once more, but now also map the variable **League** to the colour aesthetic and the variable **Salary** to the size aesthetic.

```
homeruns_plot_4 <-  
  ggplot(Hitters, aes(x = HmRun, y = Hits, colour = League, size = Salary)) +  
  geom_point() +  
  labs(x = "Hits", y = "Home runs")  
  
homeruns_plot_4
```

```
## Warning: Removed 59 rows containing missing values (geom_point).
```



5

Look at the many different geoms on the reference website.


Done.

6

Use `geom_histogram()` to create a histogram of the grades of the students in the `gg_students` dataset. Play around with the `binwidth` argument of the `geom_histogram()` function.

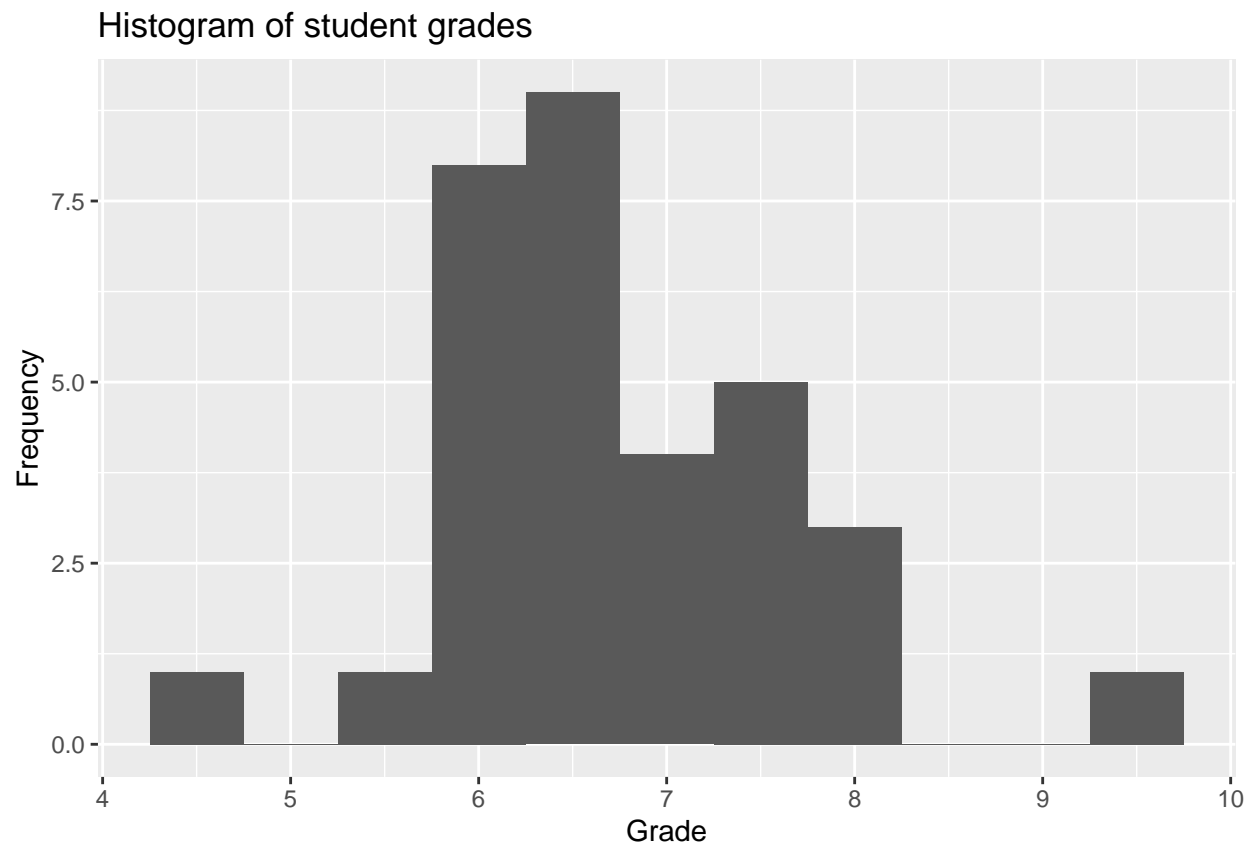
```
#binwidth = 0.1
hist_grades.1 <-
  ggplot(gg_students, aes( x = Grade))+
  geom_histogram(binwidth = 0.1) +
  labs(y = "Frequency", title = "Histogram of student grades")

hist_grades.1
```




```
#binwidth = 0.5
hist_grades.5 <-
  ggplot(gg_students, aes( x = Grade))+
  geom_histogram(binwidth = 0.5) +
  labs(y = "Frequency", title = "Histogram of student grades")

hist_grades.5
```

```
#binwidth = 1
hist_grades1 <-
  ggplot(gg_students, aes( x = Grade))+
  geom_histogram(binwidth = 1) +
  labs(y = "Frequency", title = "Histogram of student grades")

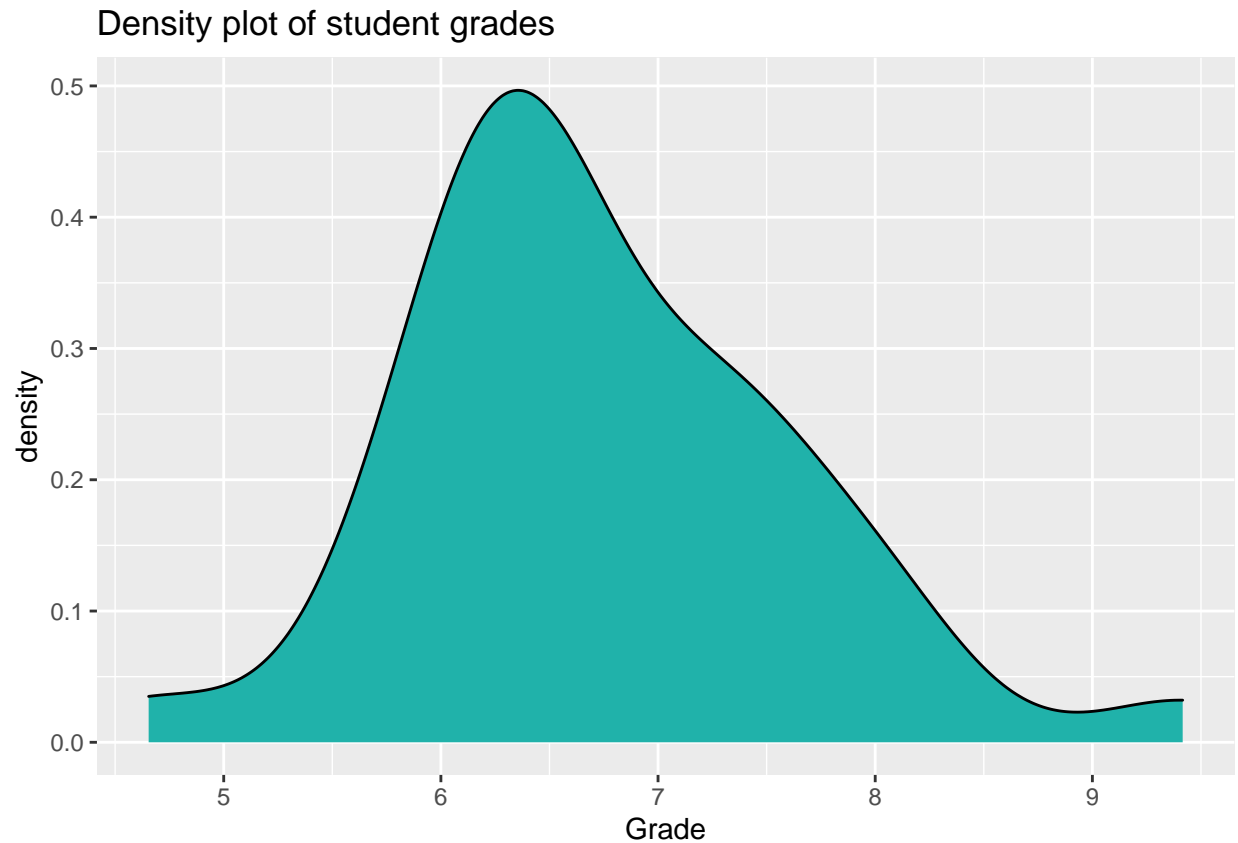
hist_grades1
```



7

Use `geom_density()` to create a density plot of the grades of the students in the `gg_students` dataset. Add the argument `fill = "light seagreen"` to `geom_density()`.


```
dens_grade <-  
  ggplot(gg_students, aes( x = Grade))+  
  geom_density(fill = "light seagreen") +  
  labs(title = "Density plot of student grades")  
  
dens_grade
```



8


Add rug marks to the density plot through `geom_rug()`. You can edit the colour and size of the rug marks using those arguments within the `geom_rug()` function.

```
dens_rug_grade <-  
  ggplot(gg_students, aes( x = Grade))+  
  geom_density(fill = "light seagreen") +  
  geom_rug(colour = "light seagreen", size = 2) +  
  labs(title = "Density plot of student grades")  
  
dens_rug_grade
```



```
dens_rug_grade_pink <-  
  ggplot(gg_students, aes( x = Grade))+  
  geom_density(fill = "light seagreen") +  
  geom_rug(colour = "deeppink2", size = 0.2) +  
  labs(title = "Density plot of student grades")
```

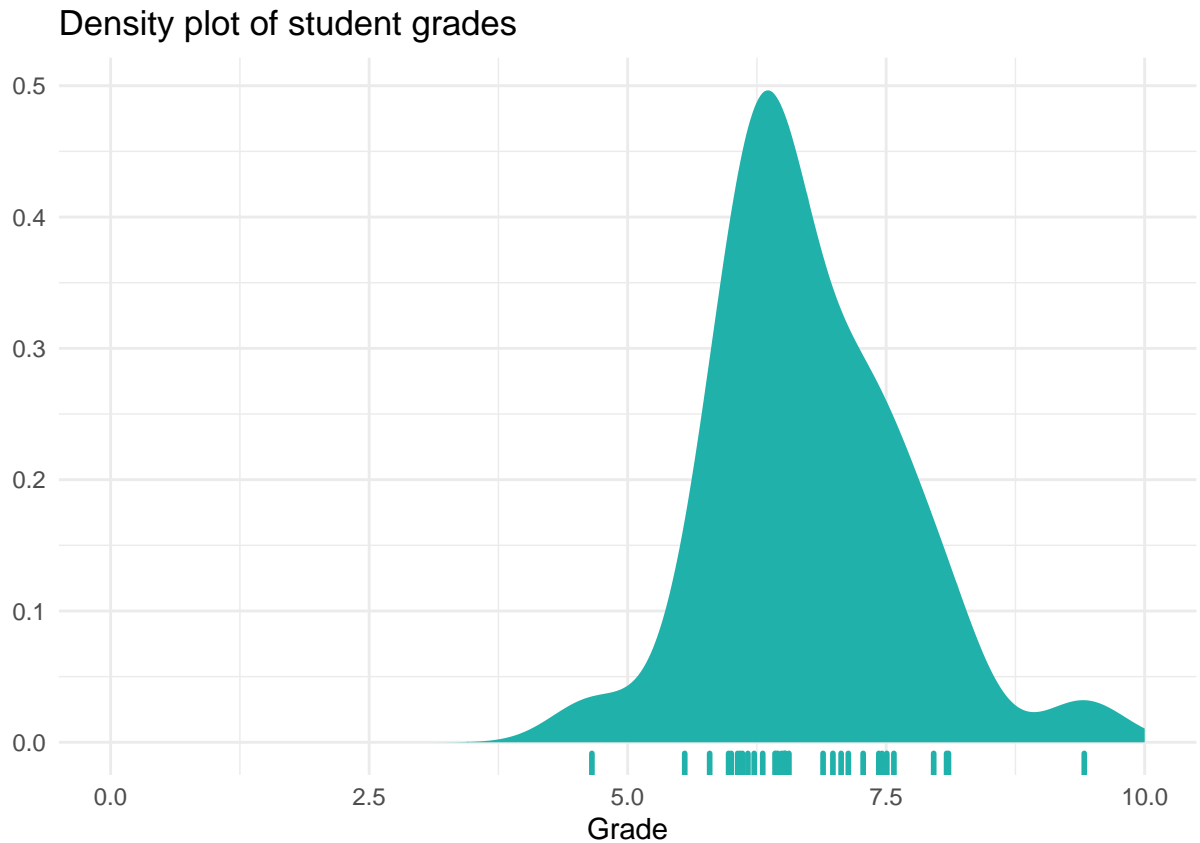
```
dens_rug_grade_pink
```



9

Increase the data to ink ratio by removing the y axis label, setting the theme to `theme_minimal()`, and removing the border of the density polygon. Also set the limits of the x-axis to go from 0 to 10 using the `xlim()` function, because those are the plausible values for a student grade.

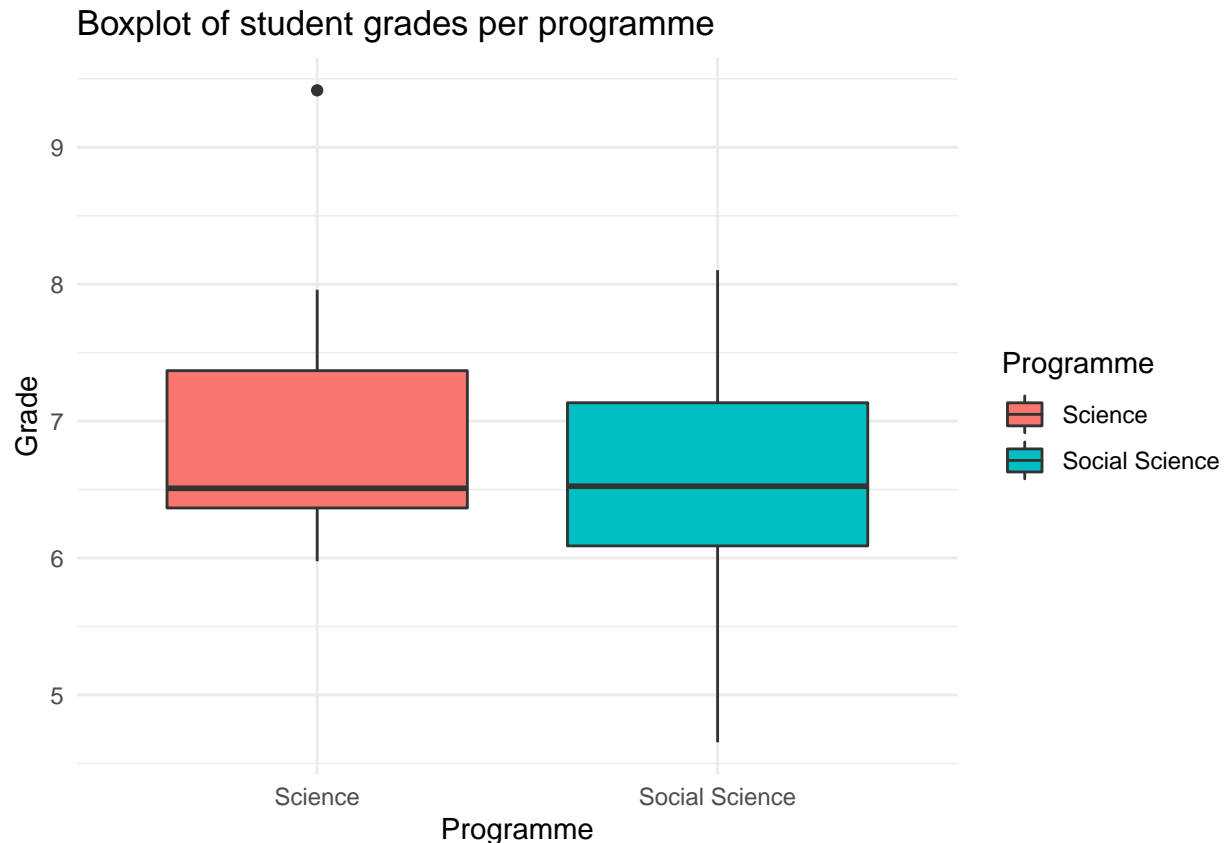
```
ink_ratio <-  
  ggplot(gg_students, aes( x = Grade))+  
  xlim(0,10) +  
  geom_rug(colour = "light seagreen", size = 1) +  
  geom_density(fill = "light seagreen", linetype = 0) +  
  labs(title = "Density plot of student grades", y = "") +  
  theme_minimal()  
  
ink_ratio
```



10

Create a boxplot of student grades per programme in the `gg_students` dataset you made earlier: map the programme variable to the x position and the grade to the y position. For extra visual aid, you can additionally map the programme variable to the fill aesthetic.

```
box_student <-  
  ggplot(gg_students, aes( x = Programme, y = Grade, fill = Programme))+  
  geom_boxplot() +  
  labs(title = "Boxplot of student grades per programme") +  
  theme_minimal()  
  
box_student
```



11

What do each of the horizontal lines in the boxplot mean? What do the vertical lines (whiskers) mean?


The horizontal lines are the first, second, and third quantiles (0.25, 0.5, and 0.75 percentile, respectively). The 0.5 quantile is also referred to as the median. The upper (lower) whiskers reaches until the largest (smallest) data point that is within 1.5 times the distance of the inter-quartile range (diff. between third and first quartiles) from the third (first) quartile. Everything beyond the whiskers is represented as dots.

12

Comparison of distributions across categories can also be done by adding a fill aesthetic to the density plot you made earlier. Try this out. To take care of the overlap, you might want to add some transparency in the `geom_density()` function using the `alpha` argument.

```
dens_transp <-  
  ggplot(gg_students, aes( x = Grade, fill = Programme))+  
  geom_density(alpha= 0.5) +  
  geom_rug(colour = "light seagreen", size = 2) +  
  labs(title = "Density plot of student grades")
```

dens_transp




13

Create a bar plot of the variable Years from the Hitters dataset.

```
years_bar <-  
  ggplot(Hitters, aes( x = Years)) +  
  geom_bar() +  
  labs(x = "Years", y = "Counts")
```


```
years_bar
```



14


Use `geom_line()` to make a line plot out of the first 200 observations of the variable `Volume` (the number of trades made on each day) of the `Smarket` dataset. You will need to create a `Day` variable using `mutate()` to map to the x-position. This variable can simply be the integers from 1 to 200. Remember, you can select the first 200 rows using `Smarket[1:200,]`.

```
vol_line <-  
  Smarket[1:200,] %>% mutate(Day= 1:200) %>%  
  ggplot(aes( x = Day, y = Volume)) +  
  geom_line() +  
  labs(x = "Day", y = "Volume")  
  
vol_line
```




```
vol_line_mod <-
  Smarket[1:200,] %>% mutate(Day= 1:200) %>%
  ggplot(aes( x = Day, y = Volume)) +
  geom_line(color = "darkred", size = 1.5) +
  geom_point(color = "darkred", size = 2) +
  labs(x = "Day", y = "Volume")

vol_line_mod
```



```
#For a good visual size= 1 is better:
vol_line_mod <-
  Smarket[1:200,] %>% mutate(Day= 1:200) %>%
  ggplot(aes( x = Day, y = Volume)) +
  geom_line(color = "darkred", size = 1) +
  geom_point(color = "darkred", size = 1) +
  labs(x = "Day", y = "Volume")

vol_line_mod
```



16

Use the function `which.max()` to find out which of the first 200 days has the highest trade volume and use the function `max()` to find out how large this volume was.

```
max_day <- which.max(Smarket[1:200,]$Volume)
max_vol <- max(Smarket[1:200,]$Volume)
Smarket[170,]$Volume
```

```
## [1] 2.33083
```

```
Smarket[which.max(Smarket[1:200,]$Volume),]$Volume
```

```
## [1] 2.33083
```


The maximal volume was measured on day 170 with a volume of 2.33083 billions.

17

Use `geom_label(aes(x = your_x, y = your_y, label = "Peak volume"))` to add a label to this day. You can use either the values or call the functions. Place the label near the peak!

```
vol_line_max <-  
  Smarket[1:200,] %>% mutate(Day= 1:200) %>%  
  ggplot(aes( x = Day, y = Volume)) +  
  ylim(0,3) +  
  geom_line() +  
  geom_label(aes(x = max_day, y = max_vol + 0.4, label = "Peak volume")) +  
  labs(x = "Day", y = "Volume")
```

vol_line_max



18

Create a data frame called `baseball` based on the `Hitters` dataset. In this data frame, create a factor variable which splits players' salary range into 3 categories. Tip: use the `filter()` function to remove the missing values, and then use the `cut()` function and assign nice labels to the categories. In addition, create a variable which indicates the proportion of career hits that was a home run.

```
baseball <- Hitters %>% filter(is.na(Salary) == FALSE)

baseball$SalaryFactor <- cut(baseball$Salary, 3, labels = c("low", "middle", "high"))

head(baseball)
```

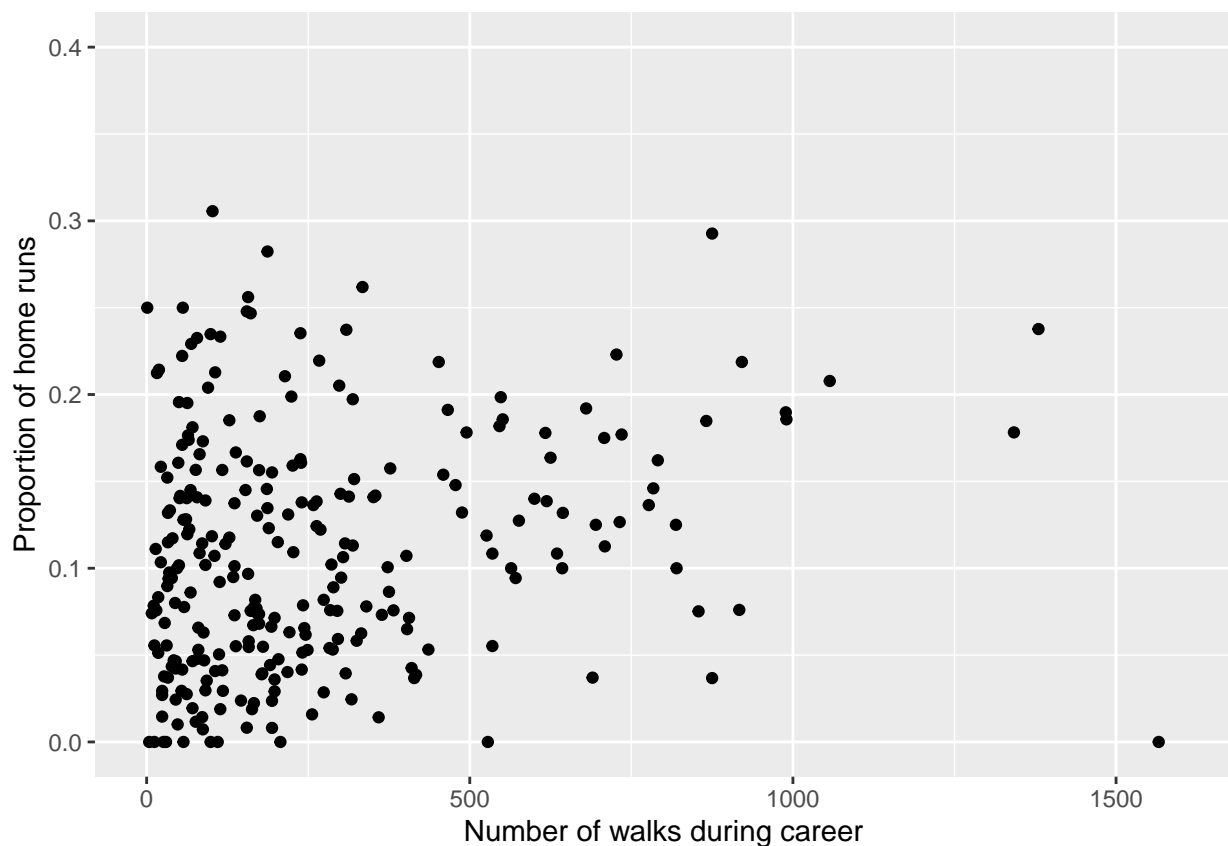
```
##           AtBat Hits HmRun Runs RBI Walks Years CATBat CHits CHmRun
## -Alan Ashby    315   81     7  24  38   39   14   3449   835    69
## -Alvin Davis   479  130    18  66  72   76    3   1624   457    63
## -Andre Dawson  496  141    20  65  78   37   11   5628  1575   225
## -Andres Galarraga 321   87    10  39  42   30    2    396   101    12
## -Alfredo Griffin 594  169     4  74  51   35   11   4408  1133    19
## -Al Newman    185   37     1  23   8   21    2    214    42     1
##           CRuns CRBI CWalks League Division PutOuts Assists Errors
## -Alan Ashby    321  414   375      N        W     632     43     10
## -Alvin Davis   224  266   263      A        W     880     82     14
## -Andre Dawson  828  838   354      N        E     200     11      3
## -Andres Galarraga 48   46    33      N        E     805     40      4
## -Alfredo Griffin 501  336   194      A        W     282    421     25
## -Al Newman     30    9    24      N        E      76    127      7
##           Salary NewLeague SalaryFactor
## -Alan Ashby    475.0      N          low
## -Alvin Davis   480.0      A          low
## -Andre Dawson  500.0      N          low
## -Andres Galarraga 91.5      N          low
## -Alfredo Griffin 750.0      A          low
## -Al Newman     70.0      A          low
```

```
baseball$Proportion <- baseball$HmRun/baseball$Hits
```

19

Create a scatter plot where you map `CWalks` to the x position and the proportion you calculated in the previous exercise to the y position. Fix the y axis limits to (0, 0.4) and the x axis to (0, 1600) using `ylim()` and `xlim()`. Add nice x and y axis titles using the `labs()` function. Save the plot as the variable `baseball_plot`.

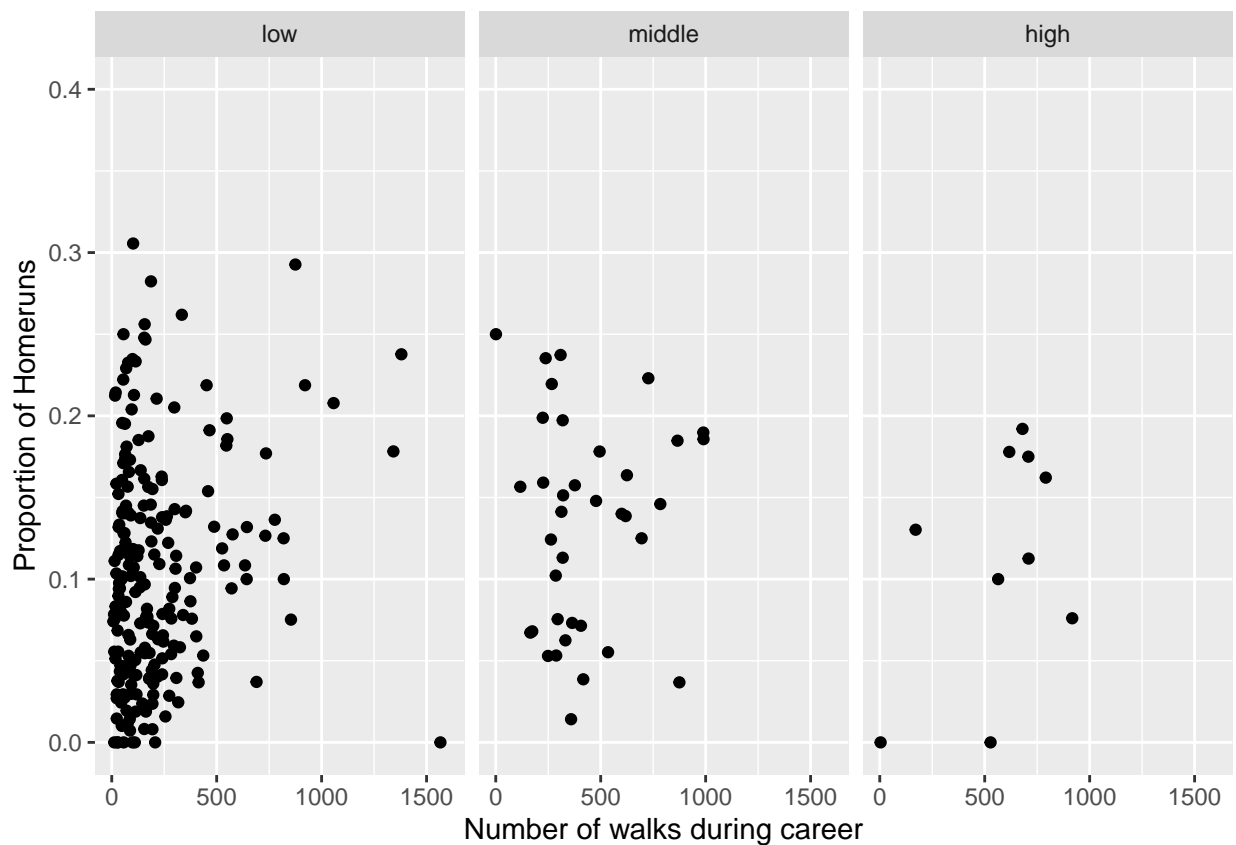
```
walks_plot <-  
  ggplot(baseball, aes(x = CWalks, y = Proportion)) +  
    ylim(0,0.4) +  
    xlim(0, 1600) +  
  
    geom_point() +  
    labs(x = "Number of walks during career", y = "Proportion of home runs")  
  
walks_plot
```



20

Split up this plot into three parts based on the salary range variable you calculated. Use the `facet_wrap()` function for this; look at the examples in the help file for tips.

```
walks_plot_split <-  
  ggplot(baseball, aes(x = CWalks, y = Proportion)) +  
    ylim(0,0.4) +  
    xlim(0, 1600) +  
    facet_wrap(vars(SalaryFactor))+  
    geom_point() +  
    labs(x = "Number of walks during career", y = "Proportion of Homeruns")  
  
walks_plot_split
```



21

Create an interesting data visualisation based on the Carseats data from the ISLR package.

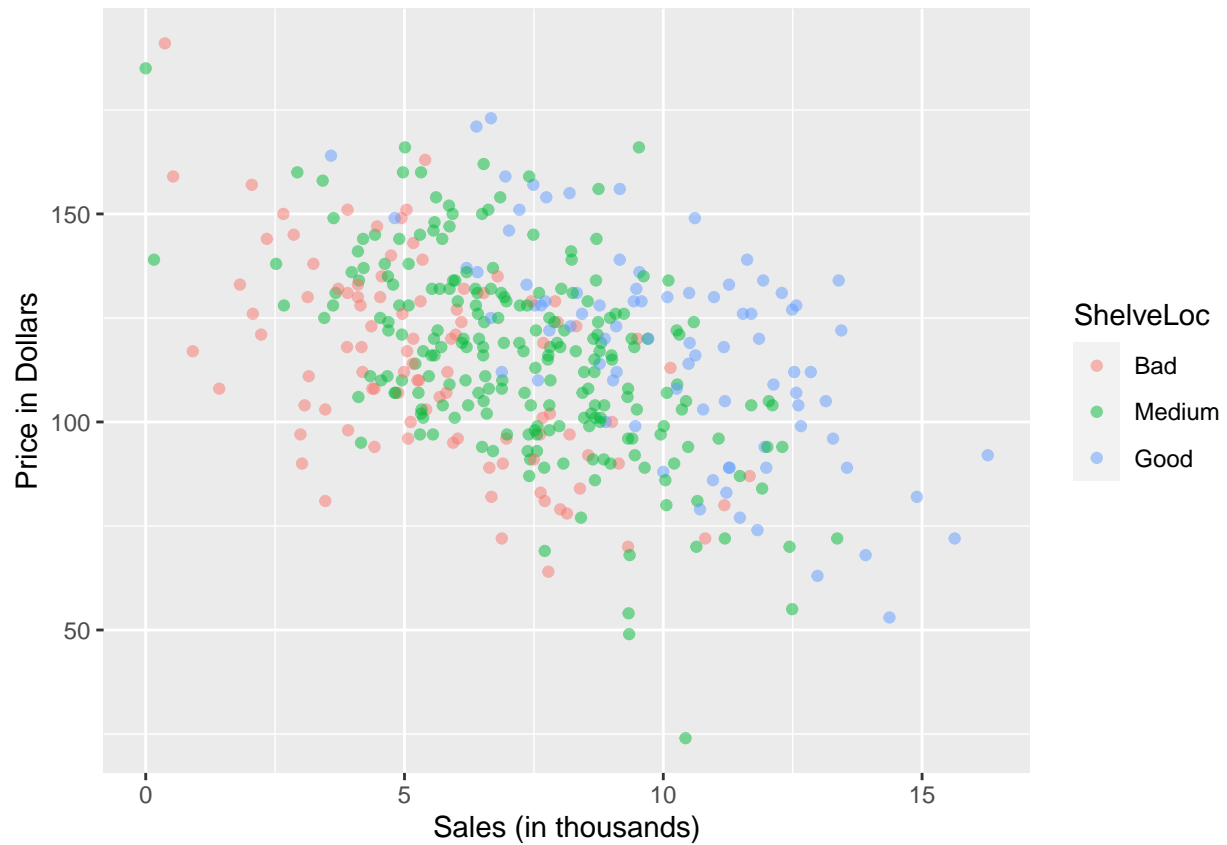
```
head(Carseats)
```

```
##   Sales CompPrice Income Advertising Population Price ShelfLoc Age Education
## 1  9.50      138     73         11         276   120      Bad   42         17
## 2 11.22      111     48         16         260    83     Good   65         10
## 3 10.06      113     35         10         269    80   Medium   59         12
## 4  7.40      117    100          4         466    97   Medium   55         14
## 5  4.15      141     64          3         340   128      Bad   38         13
## 6 10.81      124    113         13         501    72      Bad   78         16
##   Urban  US
## 1   Yes Yes
## 2   Yes Yes
## 3   Yes Yes
## 4   Yes Yes
## 5   Yes No
## 6    No Yes
```

```
Carseats$ShelveLoc <- factor(Carseats$ShelveLoc, levels = c("Bad", "Medium", "Good"))
```

```
seat_plot <-
  ggplot(Carseats, aes(x = Sales, y = Price, color = ShelfLoc)) +
  geom_point(alpha = 0.5) +
  labs(x = "Sales (in thousands)", y = "Price in Dollars")
```

```
seat_plot
```



It can be seen that (cheap) seats of good quality were sold most. Surprisingly, the most expensive seats are only of bad or medium quality, but they were also almost never sold. Another interesting observation is that there are seats in every price category from all quality categories.